

Problem 1 : Storing the dictionary

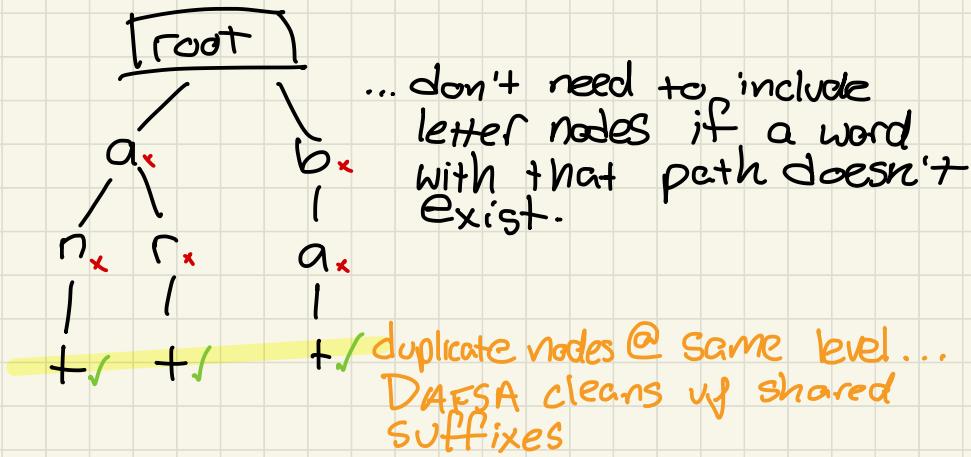
valid word bool can tell us if our ending path is a word when searching

- map?
- set?
- hash table?
- trie?

fast traversal
but large mem.
o Radix, Patricia, DAFSAs are better versions

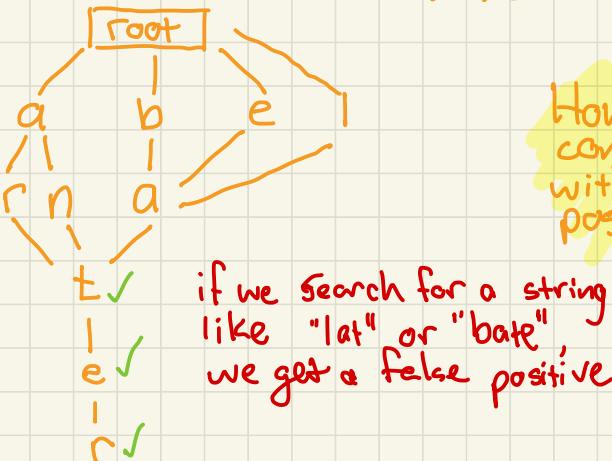
words:

ant
bat
art



words

ant
bat
art
eater
late



How do you condense vsuffixes without causing false positive searches?

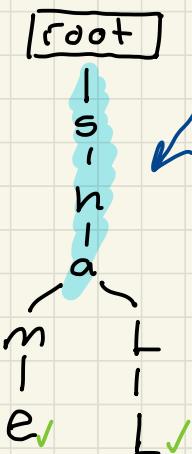


trie → LetterNode graph

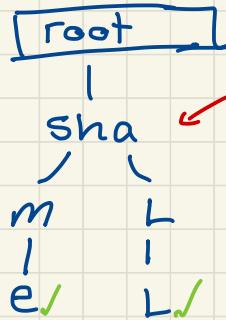
```
struct LetterNode {  
    bool isWordValid;  
    unordered_map<char, LN*> children;  
};
```

hash map
w/ const
look up time

"Jean" + rie



Patricia Trie
condense prefix nodes with only
1 parent \nRightarrow 1 child



You save on lockup
time at (maybe)
the cost of heap.

Conclusion:

- Store dictionary as a "lean" trie, a minimize function can be implemented later if heap space is an issue

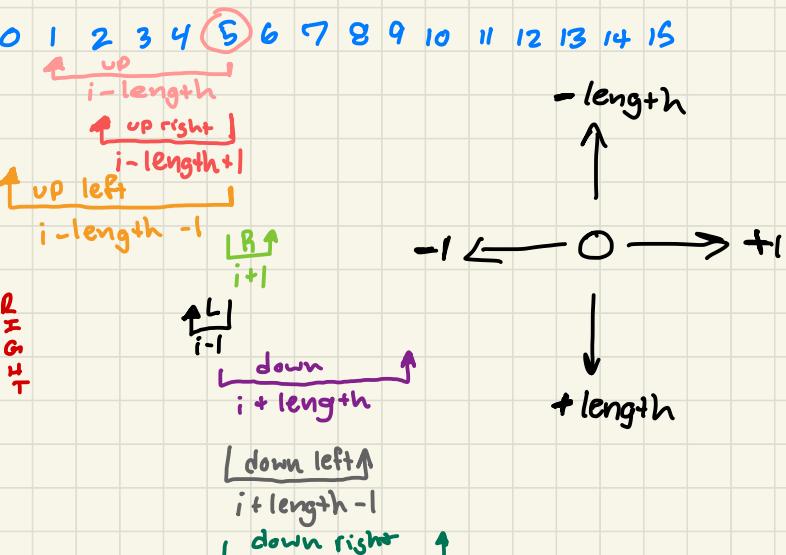
Problem 2 : Storing the board

- array or vector?

can't assume board size @ compile time.
Vector it is.

- 1d or 2d? 2d entities can be traversed in a 1d vector. Stick with 1dimension until 2d are needed.

Vector <char> 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15



Boundary Checks

$$\text{up} : (\text{index} - \text{length}) \geq 0$$

$$1 - 4 = -3$$

$$\text{down} : (\text{index} + \text{length}) < \text{vec.size()}$$

$$14 + 4 = 18$$

$$\text{right} : (\text{index} + 1) < \text{vec.size()}$$

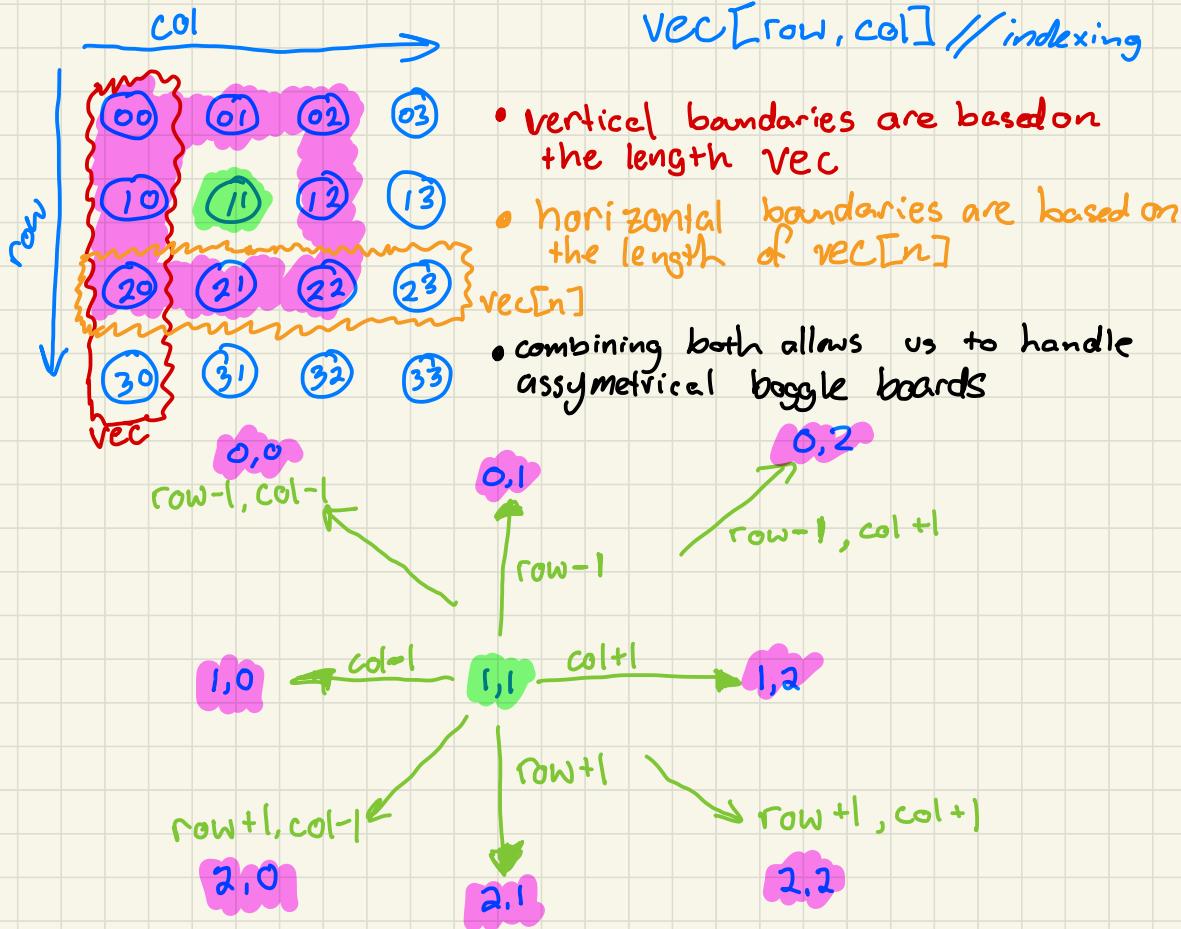
$$7 + 1 = 8 \dots$$

uh-oh... this causes word wrap



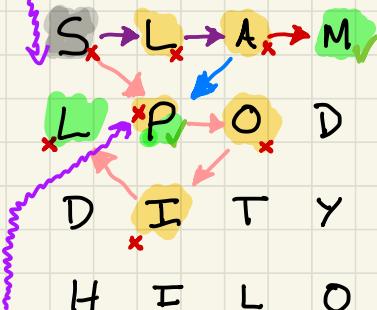
If I have to infer the boggle board line length after importing the board into a 1d vector, then I might as well just use a 2d vector with "row" "col" indexing.

New plan ... vector<vector<char>>



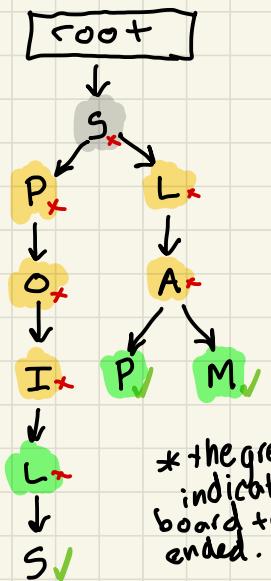
Problem 3 : board traversal alg

"SPOILS" CAN'T BE MADE B/C THE LAST "S" IS NOT VALID... ITS ALREADY A NODE IN THE PATH



Searching "sp" will not be a false positive b/c these are different 'P' nodes in the trie.

Dictionary
SLAM
SLAP
SPOILS



* the green highlight indicates where a board traversal path ended.

- pick an index to start at

- find all possible nodes
- for each of those nodes, find all possible nodes

- a valid node is a node that is not currently in a traversal path, and that forms a prefix that exists in the dictionary

- do this until you are at a node with no more options. Check if your current path is in the dictionary AND that dictionary node has its word flag set to true✓ instead of false✗

- if its a valid word, add it to an unordered_set of answers (if its not already in there... no duplicates). Since there are no more valid nodes, return and begin recursing the next node in the for each

Conclusion: this whole block is a DFS alg that can be put into a single function, given a starting index. Thread this function and give each thread a different starting index.

Problem 4: but how many threads?

- timing metrics were taken with chrono... its not the best for this kind of thing but its good enough for this exercise.

		Single Threaded Boggle Solver (Board 1 - 4x4)				Average (ms)
Debug	Trie Import Time (ms) :	25.0506	25.1517	26.193	25.1252	25.2659 25.35728
	Solver Time (ms) :	1.9846	1.8177	2.021	2.1112	2.0275 1.9924
	Total Time (ms) :	27.1762	27.1003	28.3567	27.3886	27.3977 27.4839
Release	Trie Import Time (ms) :	2.4332	2.6298	2.8343	2.585	2.6356 2.62358
	Solver Time (ms) :	0.1882	0.1927	0.2022	0.183	0.1884 0.1909
	Total Time (ms) :	2.6783	2.8864	3.1485	2.8189	2.8854 2.8835
		4 Threads Boggle Solver (Board 1 - 4x4)				Average (ms)
Debug	Thread Pool Init Time (ms) :	0.8784	0.8347	0.6762	0.834	1.1152 0.8677
	Trie Import Time (ms) :	30.545	26.8041	29.4216	27.5521	27.3681 28.33818
	Solver Time (ms) :	1.7465	1.559	1.5762	1.5344	1.5654 1.5963
Release	Total Time (ms) :	33.3179	29.2961	31.7877	30.0529	30.141 30.91912
	Thread Pool Init Time (ms) :	0.553	0.7203	0.7567	0.5431	0.6238 0.63938
	Trie Import Time (ms) :	3.4294	3.7447	4.3308	3.8123	3.6492 3.79328
Release	Solver Time (ms) :	0.0668	0.0923	0.0823	0.0992	0.0818 0.08448
	Total Time (ms) :	4.1189	4.6142	5.2264	4.5133	4.4146 4.57748
		16 Threads Boggle Solver (Board 1 - 16x16)				Average (ms)
Debug	Thread Pool Init Time (ms) :	4.2096	3.4695	3.51	2.0713	3.7923 3.41054
	Trie Import Time (ms) :	26.6416	26.3802	26.3059	28.5671	27.0124 26.98144
	Solver Time (ms) :	3.2331	3.1902	3.0763	3.2194	3.1624 3.17628
Release	Total Time (ms) :	34.2263	33.1404	32.995	34.0058	34.091 33.6917
	Thread Pool Init Time (ms) :	2.1333	3.379	1.7115	2.0407	1.5358 2.16006
	Trie Import Time (ms) :	3.5162	3.4386	3.8373	4.0508	3.6824 3.70506
Release	Solver Time (ms) :	0.1017	0.0722	0.0971	0.0865	0.1113 0.09376
	Total Time (ms) :	5.81	6.9522	5.7049	6.2431	5.3844 6.01892

Solver Time is how long it takes threads to solve the board.

- 1 thread (baseline) = 0.1909ms
- 4 threads = 0.08448ms } 4x4 boggle board
- 16 threads (hw-rec.) = 0.09376ms

4 threads 50% faster than 1 and ~10% faster than 16... maybe diminishing returns?



Threading Stress test on 16x16 board

		Single Threaded Boggle Solver (stress test - 16x16)				Average (ms)
Group D	Trie Import Time (ms) :	26.2009	27.7689	25.7078	28.2877	25.4453 26.68212
	Solver Time (ms) :	74.225	72.268	71.4945	74.2387	75.115 73.46824
	Total Time (ms) :	100.614	100.337	97.3935	102.759	100.785 100.3777
Release	Trie Import Time (ms) :	2.4715	2.4648	2.7787	2.7776	2.8215 2.66282
	Solver Time (ms) :	5.7517	5.103	5.0687	5.137	5.064 5.22488
	Total Time (ms) :	8.2816	7.6272	7.9188	7.9711	7.9414 7.94802
Group E		4 Threads Boggle Solver (stress test - 16x16)				Average (ms)
Debug	Thread Pool Init Time (ms) :	0.8697	0.7745	0.8304	0.7676	0.6836 0.78516
	Trie Import Time (ms) :	27.1962	28.5036	26.8004	30.0358	26.9143 27.89006
	Solver Time (ms) :	57.3981	55.3949	56.9946	56.6232	57.522 56.79438
Release	Total Time (ms) :	85.7466	84.8525	84.8041	87.621	85.3031 85.67368
	Thread Pool Init Time (ms) :	0.7976	0.8424	0.5025	0.6267	0.6001 0.67386
	Trie Import Time (ms) :	3.9737	3.5511	3.5012	4.3719	4.144 3.90838
Group F	Solver Time (ms) :	1.4703	1.3407	1.4524	1.2494	1.1713 1.33662
	Total Time (ms) :	6.3041	5.7963	5.5173	6.309	5.9755 5.98044
	16 Threads Boggle Solver (stress test - 16x16)				Average (ms)	
Debug	Thread Pool Init Time (ms) :	2.4742	1.9934	5.7101	3.9191	3.6894 3.55724
	Trie Import Time (ms) :	26.5714	27.3084	26.3544	27.1079	26.1164 26.6917
	Solver Time (ms) :	189.9	192.999	191.831	191.368	190.685 191.3566
Release	Total Time (ms) :	219.133	222.546	224.228	222.627	220.683 221.8434
	Thread Pool Init Time (ms) :	2.8472	2.1787	2.0709	3.3525	1.7252 2.4349
	Trie Import Time (ms) :	3.8545	4.2451	3.6	3.6127	4.1139 3.88524
Group G	Solver Time (ms) :	0.819	0.9277	0.8224	0.9376	0.9999 0.90132
	Total Time (ms) :	7.5839	7.412	6.5567	7.9758	6.9185 7.28938

- 1 thread (baseline) = 5.22488ms
- 4 thread = 1.33662ms
- 16 thread (hw-rec) = 0.90132ms

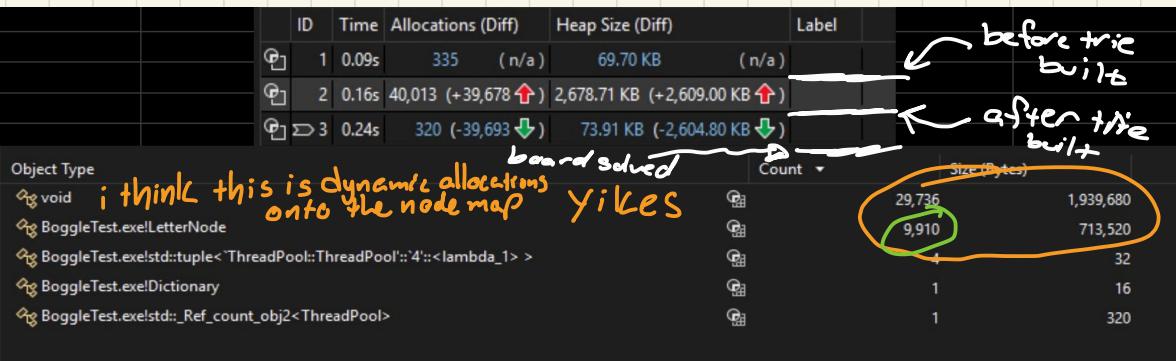
4 threads are still about 75% better than 1,
but 16 is 50% better than 4.

Conclusion :

- because varying board sizes was never specified as a requirement, default to 4 threads...
- ... but if the hw_rec is lower than 4, take the minimum of the two.
- could always dynamically decide how many threads to run based on board size, but i'm treating it as YAGNI for now.

Problem 5: how is my heap?

- memory profiles taken with VS built-in memory profiler.



A screenshot of a table showing performance metrics for a 4-threaded Boggle solver. The table has columns for Group, Configuration, Metric, and Average (ms). Red redaction marks cover most of the table content.

Group	Configuration	Metric	Average (ms)
Group B	Debug	Trie Import Time (ms) :	30.545 26.8041 29.4216 27.5521 27.3681 28.33818
	Release	Trie Import Time (ms) :	3.4294 3.7447 4.3308 3.8123 3.6492 3.79328
		Total Time (ms) :	4.1189 4.6142 5.0264 4.5133 4.4146 4.57748
Group G	Release	4 Threads Boggle Solver (Board 1 - 4x4) - Dictionary Imported Into Trie with unordered_map<char, LetterNode*>	43.95 45.27

No SPOILERS! We'll get to this data later.

- threads are now solving the board faster, but my trie import time (building the dictionary) is still a single thread. It accounts for 80% of my execution time.
- that trie process is taking up 2.6 MB, which seems like too much for a relatively small dictionary of only 9910 trie nodes
- let's convert the

```
LetterNode { map<char, LetterNode*> }
```

 to

```
LetterNode { vector<pair<char, LetterNode*>> }
```

```
LetterNode { vector<pair<char, LetterNode*>> }
```

LetterNode & vector<pair<char, LN>> heap analysis

	ID	Time	Allocations (Diff)	Heap Size (Diff)	Label	
[]	1	0.09s	324 (n/a)	68.24 KB (n/a)		
[]	2	0.13s	17,263 (+16,939 ↑)	554.55 KB (+486.31 KB ↑)		
[]	3	0.16s	320 (-16,943 ↓)	73.94 KB (-480.62 KB ↓)		
[]	4	0.16s	320 (+0)	73.94 KB (+0.00 KB)		

Object Type	Count	Size (Bytes)
BoggleTest.exe!LetterNode	9,910	317,120
void	6,986	161,632
BoggleTest.exe!std::tuple<ThreadPool::ThreadPool::'4':<lambda_1>>	4	32
BoggleTest.exe!std::_Ref_count_resource<Dictionary *,std::default_delete<Dictionary>>	1	24
BoggleTest.exe!Dictionary	1	16
BoggleTest.exe!std::_Ref_count_obj2<ThreadPool>	1	320

- Switching to a vector saves 80% of heap compared to the map.

How fast is it?

4 Threads Boggle Solver (Board 1 - 4x4) - Dictionary Imported Into Trie with unordered_map<char, LetterNode*>						Average (ms)
Debug	Trie Import Time (ms):	30.545	26.8041	29.4216	27.5521	27.3681 28.33818
Release	Trie Import Time (ms):	3.4294	3.7447	4.3308	3.8123	3.6492 3.79328
	Total Time (ms):	4.1189	4.6142	5.2264	4.5133	4.4146 4.57748
Group G						
4 Threads Boggle Solver (Board 1 - 4x4) - Dictionary Imported Into Trie with vector<pair<char, LetterNode*>>						
Debug	Trie Import Time (ms):	43.6709	43.7085	45.6689	49.6838	43.9508 45.33658
Release	Trie Import Time (ms):	2.5092	2.5212	2.4467	2.7347	2.5886 2.56008
	Total Time (ms):	3.4771	3.6239	3.6026	4.1801	3.7067 3.71808

almost there, i promise

- execution times dropped by about 30%
- it can most certainly be faster with a custom Pool allocator



LetterNodes being preallocated by a custom heap allocator

	ID	Time	Allocations (Diff)	Heap Size (Diff)	Label
①	1	0.10s	324 (n/a)	68.29 KB (n/a)	
②	2	0.12s	7,358 (+7,034 ↑)	653.26 KB (+584.97 KB ↑)	100% more than last page
③	3	0.13s	330 (-7,028 ↓)	78.27 KB (-574.99 KB ↓)	

Object Type

Object Type	Count	Size (Bytes)
void	6,988	578,814
BoggleTest.exe!std::tuple< ThreadPool::ThreadPool::'4'::<lambda_1> >	4	32
BoggleTest.exe!std::Ref_count_resource<Dictionary *, std::default_delete<Dictionary> >	1	24
BoggleTest.exe!std::Ref_count_obj2<ThreadPool>	1	320
BoggleTest.exe!Dictionary	1	64

4 Threads Boggle Solver (Board 1 - 4x4) - Dictionary Imported Into Trie with unordered_map<char, LetterNode*>						Average (ms)
Debug	Trie Import Time (ms) :	30.545	26.8041	29.4216	27.5521	27.3681 28.33818
Release	Trie Import Time (ms) :	3.4294	3.7447	4.3308	3.8123	3.6492 3.79328
	Total Time (ms) :	4.1189	4.6142	5.2264	4.5133	4.4146 4.57748
Group G 4 Threads Boggle Solver (Board 1 - 4x4) - Dictionary Imported Into Trie with vector<pair<char, LetterNode*>>						
Debug	Trie Import Time (ms) :	43.6709	43.7085	45.6689	49.6838	43.9508 45.33658
Release	Trie Import Time (ms) :	2.5092	2.5212	2.4467	2.7347	2.5886 2.56008
	Total Time (ms) :	3.4771	3.6239	3.6026	4.1801	3.7067 3.71808
Group H 4 Threads Boggle Solver (Board 1 - 4x4) - Dictionary Imported with Custom Allocator (10,000 nodes)						
Debug	Trie Import Time (ms) :	46.3492	44.267	42.135	46.6357	47.6528 45.40704
Release	Trie Import Time (ms) :	2.2395	1.8793	1.9009	2.0749	2.2506 2.06904
	Total Time (ms) :	3.3356	3.1979	3.6367	3.2887	3.5256 3.3969

- heap usage is slightly larger by about 25% (100KB)
- BUT its 20% faster than the vector trie w/o the pool allocator, AND its a total 40ish% faster than where we started (trie with maps, no custom allocator)

Conclusion: faster execution time is worth the extra heap usage. This time likely is gained by not using any "new" keywords as we import the dictionary.

Note: this data uses a pool of 10K nodes, but I set the deliverable to 15K b/c idk if this will be tested with a larger dictionary. The deliverable exits gracefully b/c logs error if pool is empty.